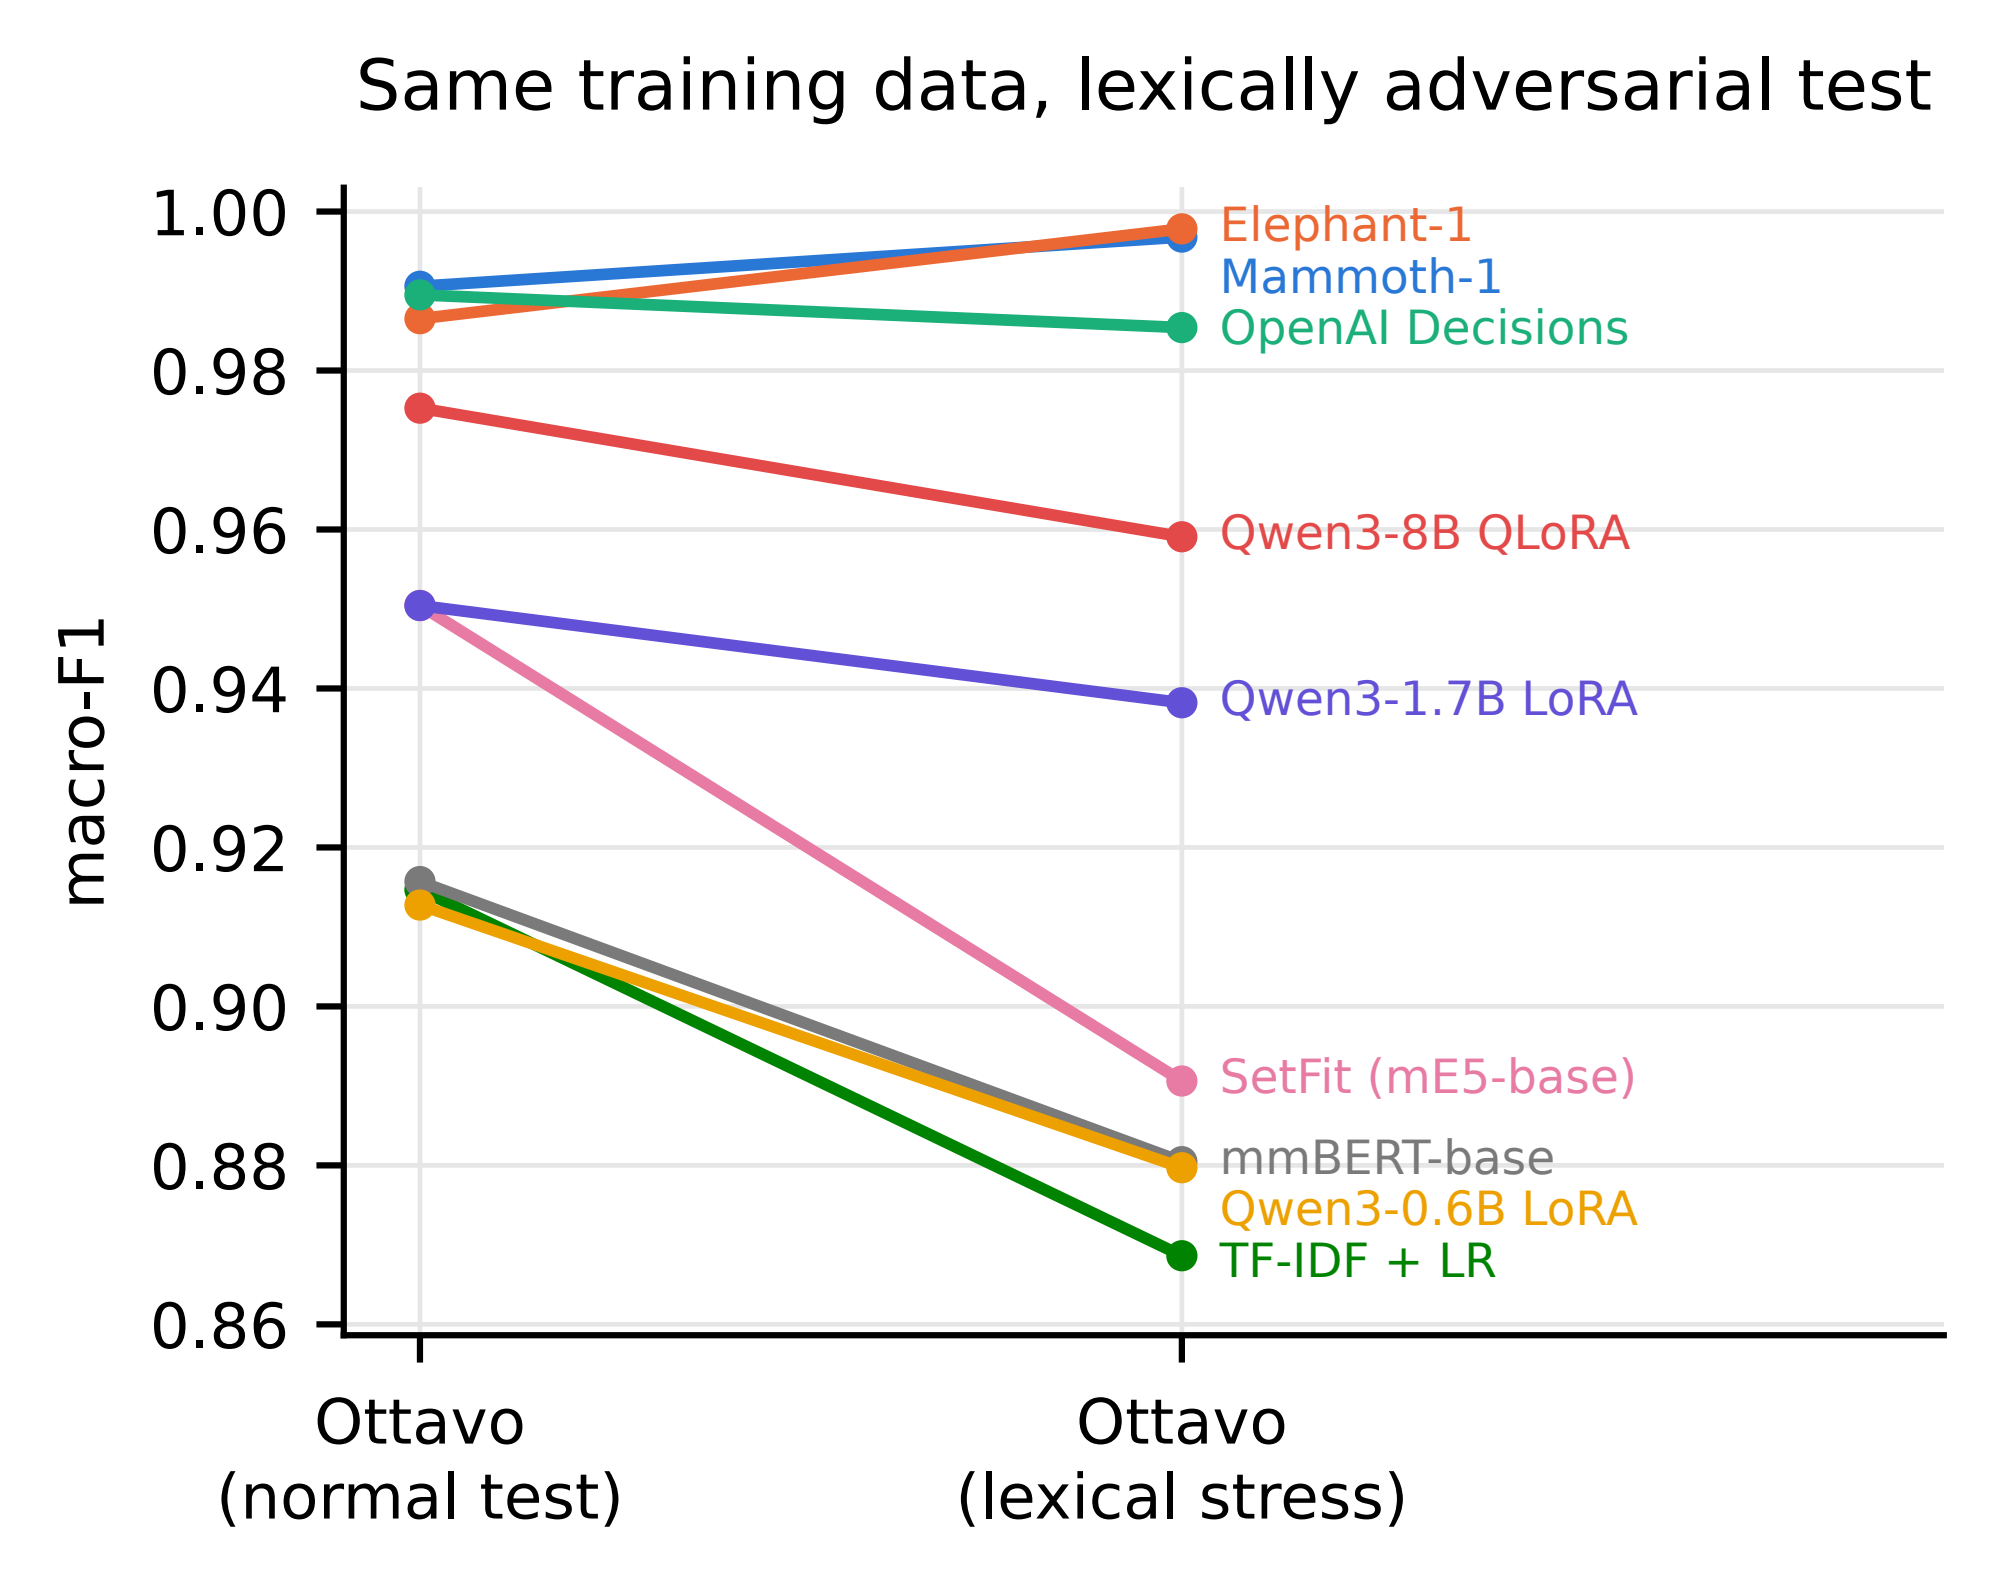
Same training data, lexically adversarial test
macro-F1
1.00
0.98
0.96
0.94
0.92
0.90
0.88
0.86
Ottavo
(normal test)
Ottavo
(lexical stress)
Elephant-1
Mammoth-1
OpenAI Decisions
Qwen3-8B QLoRA
Qwen3-1.7B LoRA
SetFit (mE5-base)
mmBERT-base
Qwen3-0.6B LoRA
TF-IDF + LR